

Dr. Alex Rivera

Senior Machine Learning Engineer | NLP Specialist

Email: alex.rivera.ai@example.com | GitHub: github.com/alexrivera-ml | Location: San Francisco, CA

PROFESSIONAL SUMMARY:

Senior AI/ML Engineer with 6+ years of experience architecting and deploying production NLP and deep learning systems. Proven expertise in Natural Language Processing, Transformer models, PyTorch, TensorFlow, Scikit-Learn, and low-latency FastAPI model serving. Experienced in MLOps, Docker containerization, Kubernetes orchestration, and AWS cloud pipelines.

TECHNICAL SKILLS:

- Languages & Frameworks: Python, PyTorch, TensorFlow, Scikit-Learn, Hugging Face, NumPy, Pandas
- NLP & Information Retrieval: TF-IDF, Word Embeddings, Transformers, BERT, LLM Fine-tuning, Tokenization, Text Preprocessing
- Backend & APIs: FastAPI, Flask, REST APIs, Pydantic, Uvicorn, Asynchronous Programming
- Cloud & MLOps: AWS (S3, EC2, SageMaker), Docker, Kubernetes, CI/CD, MLflow, Git
- Databases & Tools: PostgreSQL, SQL, Vector Databases, Pytest, Linux, Agile/Scrum

WORK EXPERIENCE:

Senior Machine Learning Engineer | DeepData Labs (2021 – Present)

- Designed and scaled transformer-based NLP pipelines serving 10M+ daily inference queries.
- Built low-latency REST APIs using Python and FastAPI for real-time document classification and ranking.
- Implemented automated CI/CD and Dockerized deployment workflows on AWS Kubernetes clusters.
- Fine-tuned open-source LLMs and built TF-IDF baseline search systems with high precision and recall.

Machine Learning Engineer | Apex Analytics (2018 – 2021)

- Developed text preprocessing, tokenization, and vectorization modules using Scikit-Learn and NLTK.
- Deployed regression and classification models using Python, Docker, and PostgreSQL.
- Conducted unit testing with pytest and maintained 90%+ code coverage.

EDUCATION:

- Ph.D. in Computer Science (Machine Learning Focus), Stanford University (2018)
- B.S. in Computer Science & Applied Mathematics, UC Berkeley (2014)